

Side effect minimization in Reinforcement Learning

Jonatan Hellgren
under supervision of: Olle Häggström

April 2022

Contents

1	Introduction	1
2	Background	2
2.1	Artificial intelligence	2
2.1.1	Future progress	4
2.1.2	Basic drives	5
2.1.3	Timeline for transformative AI breakthrough	6
2.2	AI Safety	7
2.2.1	AI alignment	7
2.2.2	Consequences with unaligned AI	8
2.2.3	Problems in AI safety	9
2.3	Approaches for creating safe AI	9
2.3.1	Learning human intent as a priority	10
2.3.2	Implementing interruptibility and corrigibility	10
2.3.3	Side effect minimization	10
2.4	Aim of thesis	11
3	Theoretical background	12
3.1	Preliminaries	12
3.1.1	Markov decision process	12
3.1.2	Deep learning	13
3.2	Reinforcement learning	13
3.2.1	Q-learning	13
3.2.2	Policy gradient	14
3.3	Side effect minimization using value-difference measures	14
3.3.1	Baselines	14
3.3.2	Deviation measures	15
4	Methods	18
4.1	SafeLife	18
5	Results	19

6	Discussion	20
7	Conclusion	21

1: Introduction

Since having a clear understanding of what will be covered in this report is crucial, we are going to begin with defining it. What will be covered in this report will be a small part of the big and hard problem of creating safe Artificial Intelligence. This is a problem of great importance that we should not overlook, because the consequences of what we manifest in the present or near future, may last for our remaining history.

We will begin by defining the concept. Then go through how it works today, where current progress leads, and how fast it may be. After that we will cover the topic of why this progress can lead to terrific consequences and how we are trying to mitigate these.

1.1 Artificial intelligence

In recent human history, we have seen massive technological development. Today our lives are in several ways different compared to centuries ago. Most of this development can be seen as a consequence of the development in tools. In early prehistory, these tools were things such as fire to cook our food or spears and knives to hunt with. Later in our history, we can see that these tools tend towards more complexity in things such as mechanical machines and the printing press. In recent years a new tool has emerged and is currently starting to show its potential, namely artificial intelligence (AI).

The definitions of AI varies, likely due to the largeness of the field. On Wikipedia we find the definition[**Wikipedia**]:

Artificial Intelligence (AI) is intelligence demonstrated by machines, as opposed to the natural intelligence displayed by animals including humans.

However to understand this definition properly it is necessary to also define what intelligence is. In [**Tnkande Maskiner**] the author brings up the following two definitions to clarify this: “the quality that enables an entity to function effectively and with foresight in its environment” and “the ability to correctly perceive one’s surrounding environment and act in away that maximizes one’s chances of achieving given goals”.

In the standard textbook on the subject [**Russel Norvig**] the authors define AI with a suitingly wide definition. They state that the field of AI is, “concerned with not

just understanding but also building intelligent entities - machines that can compute how to act effectively and safely in a wide range of novel situations". They later goes on describing four different approaches the field can be categorized in to. The approaches involve making an AI that can either act or think in a way that is humanly or rational. Where rationally is considered as a more abstract and formal definition of intelligence, that basically means doing the right thing.

Ordinary computer programs is written in a code containing step-by-step instructions that a can computer execute in order to perform the desired task. This was also the case for the first two paradigms of AI: rule-based AI and expert systems[K'lla]. In these paradigms human knowledge where explicitly programmed in to the computer in order to create automation. The types of models created in such a way is typically good for less complex goals where everything can be explicitly modelled. However, in the current paradigm of machine learning the task is to create a model that can process information faster and recall a greater quantity then humans[k'lla]. This allows for automation in more complex task where explicitly defining everything is infeasible. L'S CHAP 1 Russel Norvig

AI has in recent years been applied in the industry more broadly and it is already generating a yearly revenue of trillions of dollars[Russel Norvig]. This is mostly due to the recent and impressive progress in the current paradigm. This progress has in recent years become a possibility due to more data being available, faster computer hardware, and the massive amount of funding that is spent on research. Although these systems are quite automated, a key point here is that these systems still require humans to create and function.

In [Russel Norvig] they call the path of creating AI systems that act rationally "The rational agent approach". An agent is something that acts or more specifically is able to: operate autonomously, perceive the environment, persist over a prolonged time period, adapt to change, and create and pursue goals.

The development of AI agents shifts it from a tool to a autonomous tool. We have already seen this shift happen in many situations[k'lla]. The reasoning why this is attractive is that the human intervention part required by an AI tool is likely to become a bottleneck[T'nkande Maskiner]. Since, human intervention is likely to become a bottleneck in both intelligence and speed.

The biggest field of creating intelligent agents is called Reinforcement Learning (RL). RL is similar to how one goes about to train a pet, with desirable behaviour a positive rewards is given while undesired behaviour is discouraged with negative rewards. This field has seen a substantial development in the recent years with advances in board games such as chess and go[Silver et al.], autonomous vehicles[Levinson et al.], and video games[Minh et al.]. These advancements motivates the usefulness of implementing such agents more broadly in our daily life.

1.1.1 Future progress

When the pioneers in the field of AI started the development, the ideas were not to apply systems that automate a narrow set of tasks, as we can see in modern AI systems. The ideal was instead to recreate the intellect of a human in a machine[**McCarthy et al.**]. To extend our thoughts from mere thoughts to a new life form with a base of silicon-based hardware instead of carbon-based wetware. This is often referred to as Artificial General Intelligence (AGI), which is an AI that can solve an arbitrary set of tasks with as good or better performance than a human is capable of. The main difference from AI is that the set of tasks is not bounded.

Take for example DeepMinds AI system AlphaGo that won against the world champion Lee Sedol in the game of Go [**DeepMind**], if we were to apply the same system on the task of sorting mail, it would fail spectacularly. The reason is the team of brilliant researchers at DeepMind designed the model specifically to be good at Go¹. An AGI would on the other hand be able to play a game of Go, then drive its car, to do its job where it sorts mail and much more.

The significant difference with this shift is that it will increase the possible tasks that a single system can perform. The possible tasks would become arbitrary and be performed at a human level or higher. The implications of such a breakthrough would likely be on the same scale as the industrial revolution[**Critch Kruger**], but instead of automating physical labor we would instead have automated mental labor. The following quote summarizes the potential impacts “Machine intelligence is the last invention that humanity need ever to make”[**I.J Good**]. This could be understood by realizing that for every possible invention we could come up with and every possible labor, the machine would be able to either invent or automate.

A reason to believe that such systems are possible to build, is that we know the human intelligence was able to evolve naturally with evolution, thus something similar should be possible to reproduce in machines. There are arguments that say the created intelligence could more intelligent then us, since intelligence might not have been selected for by evolution[**S. Legg**]. When we develop AI we can focus the development for specifically on intelligence. This is as long as we do not believe in substance dependence [**Bostrom (2003)**], that is to believe that intelligence can only occur in carbon-based life forms and not in silicon-based.

Although it has been argued that an AGI breakthrough is not necessary to have a large impact on our world, because a lot of things we humans deem as intelligent will not help the AI in doing so. Take for example speech, if an AI could create convincing and motivating speeches, then it would be enough for having a large effect on politics and thus impact legislation and policy making. Another one is finance, where a potential AI could steer the world’s funding towards its specific goals. For this rea-

¹In more recent years DeepMind has released a new AI called AlphaZero which has a more general approach and is thus able to play Go, Chess, and Shogi[**Deepmind2**]. Nevertheless, the set of tasks is still limited. A finite two-player zero-sum board game.

son, many researchers have stopped talking about AGI, and have instead refined the concepts[**Critch Kruger**]. An AI system that is capable enough to induce transformative consequences on the same scale as the industrial or agricultural revolution is called an *transformative AI* (TAI). On the other hand, if this *transformative AI* also would be unstoppable once deployed it is called an *prepotent AI*.

Developing a TAI is not an easy task, however it might not be necessary to create one directly in order for it to be created[**Superintelligence**]. A different approach is to create a AI system that can develop an TAI system. A key property for this AI system is self improvement. Theoretically if an AI system has reached threshold where it is better at improving it self better then its creators. By letting an AI create the next version of it self the new version would become even at better self improvement. If this iterative process keeps going it would create a intelligence explosion[**Yudkowsky**] often referred to as the singularity.

1.1.2 Basic drives

Understanding what impacts a potential TAI or AGI will cause is hard without understanding how it will behave. There has been a lot of work laying the foundations for understanding the possible behavior by hypothesising about what drives it could have. A commonly adopted view (but still controversial[**Miller Cannon**]) is the Omohundro-Bostrom theory for AI driving forces. Two cornerstones together imply it[**O Hggstom**], namely *instrumental convergence thesis* and the *orthogonality thesis*, which we will now explain further.

The AI systems of today typically are applied by giving an agent a goal, called the terminal goal. This goal could be anything, for example maximizing the number of paper clips produced by a factory, solving the Riemann hypothesis, or counting all the blades of grass on our planet. When the system pursuits its terminal goal, it is rewarded.

When a agent pursuits this goal, there would naturally arise other instrumental goals. Examples of such would be self-preservation, self-improvement, discretization, goal perseverance, and resource accumulation[**Omohundro**]. The reasoning behind this is that these instrumental goals help the agent in the pursuit of its terminal goal. The agent wouldn't be able to perform its goal if it were destroyed for example and thus self-preservation would arise.

These instrumental goals will likely be shared between a wide range of different agents, since pursuing them helps the agent achieving its terminal goal, regardless of what it is. Thus there is a set of instrumental goals which agents would naturally converge towards and hence the name.

To this day there does not yet exist any rigorous mathematical proof for this. Some work has however been done in trying to lay the necessary foundations for it [TURNER et al]. In the paper, the authors prove in a simple environment that certain actions gives the agent more power in the sense that a wider range of possible future actions become available. On average it is optimal to choose those actions that yields

higher power. Thus we can see the pursuit of instrumental goals as tendency to seek power.

The orthogonality thesis was first described by Nick Bostrom[**Bostrom2**], it states that the intelligence of an AI is logically independent of the goals it might have. Thus a very intelligent AI could in theory have from our point of view a stupid task, such as counting all the blades of grass on our planet. Or it can have a goal that we may deem as an important one, like keeping the climate on earth habitable for the species that currently live on it. For an AI both of these tasks would be as important, given that we assigned the goal to it during its creation. The same would be the case for a not-so-intelligent AI.

1.1.3 Timeline for transformative AI breakthrough

The well known and impressive AIs of today still have not reached the levels required for an TAI. For example AlphaStar a AI that won against world champions in the complex computer game DotA 2[**Deepmind**] is estimated to “about as sophisticated” as a bee[**A Cotra**]. While the state-of-the-art language generator model that can summarize, continue, and carry out convincing conversations is estimated to be “more sophisticated” than a bee[**A Cotra**].

This raises the question: When we will see these breakthroughs in the field that enables the creations of TAI systems? This is hard to answer, but with all the focus in the form of funding[**K’lla**] and research[**k’lla**] that is applied to it, we are undoubtedly getting ever closer. There have been some research on the matter and the results of a survey and a more quantitative forecasting model will in this subsection be presented.

In a well cited survey [**Grace et al**] (2017) they asked researchers in the field of AI to estimate the probability of human-level machine intelligence (unaided machines that can achieve all tasks better and more cheaply than human workers) arriving in the future years. The conclusion of the survey where:

Researchers believe there is a 50% chance of AI outperforming humans in all tasks in 45 years and of automating all human jobs in 120 years, with Asian respondents expecting these dates much sooner than North Americans.

Although this result should be taken with a grain of salt since the distribution of answers had a large variation. Also, seemingly there should not be such a big difference between solving all task and all jobs, since a jobs consists of a set of task and thus if one can perform every task one should be able to perform every job. There is still something we can take away from this survey about the timeline, namely that researches mostly thinks all tasks will be automated within this century. But, perhaps it says more about how unsure the research field is.

In the quantitative forecasting model by [**Ajeya Cotra**], they present a model that predicts when we will be able to train an TAI system. This study uses biological anchors

in order to estimate how much compute is necessary for the training. These anchors are based on factors that played a role for the development of human intelligence. The sizes of these factors range from the amount of information in our genome, to the computational power in our brain, and all the computational power available on our planet. Each anchor is weighed according to how likely the author believes them to be.

Then using parameters such as rate of development in hardware, algorithmic progress, and willingness to spend money, they are able to estimate how likely it is that an TAI system will be developed for any given year in the future.

The results were summarized by [Robin Shah AN“#121] as the following:

For the median of 2052, the author guesses that these considerations roughly cancel out, and so rounds the median for development of TAI to 2050. A sensitivity analysis concludes that 2040 is the “most aggressive plausible median”, while the “most conservative plausible median” is 2080.

This forecast presents a shorter timeline compared to the previously presented survey, but it also answers a different question so they cannot be compared directly. Although together they agree that we will likely see the development of TAI systems this century.

There is one thing worth mentioning when talking about the timelines for future TAI. It is not necessarily true that the amount of progress will continue to develop at the current rate, it could either decrease or increase. The field of AI has previously been through two winters where the funding and excitement decreased[Russel Norvig]. This was mainly due to high expectations that were not met. So if a third winter were to emerge we could expect the rate of development to decrease. Contrary, the progress could significantly increase due to breakthroughs in relevant fields and thus shorten the timeline.

1.2 AI Safety

All tools can be applied in multiple ways, some might be beneficial and some might be ill-intentioned. Take for example a hammer, you could either use it to build a house where you can live with your family, or you could use it to hit another person. The same is the case for AI because it still is but a tool. Although, the consequences might be more severe and possibly even existential[K’lla]. Also, we can not guarantee that even a well intentioned use will be safe, this will be covered later in this section.

We will now take a closer look at what we define as a safe AI and what problems could arise if we fail to make it.

1.2.1 AI alignment

- AI alignment as an approach for safe AI

- Clarifying alignment, inner and outer
- Explain figure




Figure 1.1

1.2.2 Consequences with unaligned AI

The task of creating safe AI is hard, mainly since humans evolved to understand other humans, not computers. Elizer Yudkowsky explained that this becomes a problem “because it will be able to find solutions we can not think about” [**Yudkowsky speech**], since it can look for solutions in a completely different and eventually larger solution space.

In [**Critch Kruger**] they present the human fragility argument, it attempts to clearly explain why unaligned AI in the future could become a existential threat to humanity. It states:

The human fragility argument. Most potential future states of the Earth are unsurvivable to humanity. Therefore, deploying a prepotent AI system absent any effort to render it safe to humanity is likely to realize a future state which is unsurvivable.

The first part of it can be understood by for example realizing that we are fragile to changes in the atmosphere, temperature, and our ecosystem. Since a prepotent AI by definition will make large impact and be unstoppable ones turned on, we can not guarantee that the changes made wont affect the things we are fragile towards.

In the upcoming century Toby Ord, a philosopher that focuses on existential risk, loosely estimates that the chance of humanity facing an existential catastrophe is 1 in 6, out of which a chance of 1 in 10 are due to unaligned AI [**precipice**]. He arrived at this conclusion by estimating a 50% chance for a prepotent AI breakthrough and a 20% chance of failure with the alignment of that system [**rationaly speaking**].

It is however necessary to point out that this is only an estimate that is meant to express the importance of the problem and should not be taken as a fact. The key takeaway here is that there is a quite large chance of facing an existential threat due to future unaligned AI. Also that he believes that unaligned AI poses the highest probability for existential risk in the upcoming century, where other causes were things such as an asteroid impact, nuclear war, and pandemics.

1.2.3 Problems in AI safety

Reward functions are very hard to specify[Turner et al. (2020)], such that they can not be exploited by an agent once employed. Exploiting here refers to when a behavior is developed by the agent that optimizes the reward without performing the task it was meant to learn as intended. This is called *reward hacking*.

A real life example of reward hacking include is; When training a robotic vacuum cleaner to drive more carefully and not bump into things hard, by yielding a negative reward based on how hard it bumped in to obstacles. It developed a behaviour that instead of driving slowly when approaching obstacles, to drive backwards since there were no bumpers on the back and thus no negative reward[Custard Smingleigh].

If we want to measure if the robot is cleaning cautiously, then measuring the force that the bumper senses is a good measure. But, when letting it create a behavior that minimizes this unwanted side effects may arise. This can be seen as a consequence of Goodhart’s law, which states that: “When a measure becomes a target, it ceases to be a good measure”[Goodhars-wiki].

In addition to the difficulty of specifying a proper reward function, negative side effects may also arise as a unintended consequence of a proper optimal behaviour. In [Saisubramanian et al] they state that negative side effect “occur because the agent’s model and objective function focus on some aspects of the environment but its operation could impact additional aspects of the environment”.

These problems are alarming since if we have a problem with the AI of today, how severe might future problems be with more powerful AIs that also might be applied more broadly. Several AI researchers have raised warnings for future development of AI, Stuart Russel, Max Tegmark, Eliezer Yudkowsky to name but a few. [K”LLOR]

A common and rather cartoonish example of how it can go wrong is the paperclip armageddon described in *Superintelligence*. In it, there is a paperclip factory that has an AI which maximizes the amounts of paperclips created in the factory. In an update, the system is accidentally transitioned to the level of an AGI. Eventually, the paperclip maximizer comes to a point where the existence of humans serves no purpose or possibly even negatively affects producing paperclips, and thus they become extinct.

This examples illustrates that a seemingly stupid task can be seen as more important to an AI than the existence of the human race on the planet, if we where to program it as its goal. Another is that a goal given to an AI does not need to sound harmful in order to pose an existential risk.

1.3 Approaches for creating safe AI

The research field of creating safe and aligned AI has in recent years seen a substantial increase. We are however a long way from solving the problem, most of what is being done today are mainly speculations and laying necessary foundations for future research.

Solving this issue in time is extremely important since if we see the emergence of a transformative AI or possibly even an unstoppable prepotent AI, humanity might suffer the consequences previously described.

There are several proposed paths for solving this issue and perhaps the sheer amount might signify the difficulty of the problem. We will now take a closer look at a few of these paths in this section.

1.3.1 Learning human intent as a priority

- Inverse Reinforcement Learning, what it is and key ideas behind
- Solving outer alignment by making agent unsure of what human intent is
- Potential issue currently

1.3.2 Implementing interruptibility and corrigibility

- Why not turn it off when it goes badly?
- Allowing modifications of objective function and hitting off switch

1.3.3 Side effect minimization

Attempts have been made to limit these side effects by specifically specifying what the agent should not do[Zhang et al]. However, with this approach, the creation of the reward function becomes an iterative trial and error process. This requires a lot of human intervention, which makes the agent less autonomous and requires more time.

To solve this, attempts have been made to define a set of constraints that makes the agent avoid side effects without the need to specify what a side effect is. It is also important that the constraints defined should be able to extrapolate into new unseen situations.

An example of was presented in [Armstrong and Levinstein], where they measured the impact as the difference in the world if the agent were turned on compared to if it was not turned on, where the world is simplified as a set of parameters. However, the choice of parameters will either be quite large or chosen quite arbitrary. But this idea laid the philosophical groundwork for future solutions.

A more general approach to defining side effects is presented in [Eysenbach et al], where the agent is penalized if they are not able to preserve reachability to the initial or any other defined safe state. This method incentivizes a safe exploration that avoids irreversible states. This works well when no such irreversible action is required and the agent to reach its goal, to make an omelet one has to break some eggs. Another problem arises when the agent is in a dynamic environment, since then it would act to prevent other irreversible actions from happening, like a human eating an omelet.

In [Krakovna et al 2019] and [Turner et al 2020], the method for defining side effects is done by defining a baseline, and a deviation measure from that baseline. This approach allows for different baselines and deviation measures, thus it allows for an even more general approach. A baseline can be the initial state the agent where deployed in or as previously brought up, if the agent was not turned on. The deviation measure can be things such as how well the agent preserves possibilities for reaching future states or performing auxiliary tasks.

This type of method is what this report will focus on. The details of this method will be further explained in the theory chapter, once some necessary preliminaries have been covered.

1.4 Aim of thesis

This thesis aims to investigate how variations of current methods that reduce side effects by including a value difference measurement compare to standard methods.

2: Theoretical background

Since we have not yet reached an TAI or AGI breakthrough yet it is not possible to test methods of side effect minimizations on them directly. Instead we have to make use of what we currently have available for us. Recent promising results in RL motivates the use for it as a substitute for creating intelligent agents.

In this chapter we will cover the basics of RL and two approaches for creating them. But, before that we will take a look at some preliminary theory that is used in RL. After this we will describe methods for side effect minimization in more detail.

2.1 Preliminaries

2.1.1 Markov decision process

A Markov decision process is a stochastic decision process, where an agent is navigating it. The Markov property implies that the process is memoryless, meaning that the previous state do not have an effect on the next choice, only the current one does. In mathematical terms it can be described as,

$$p(a_t | s_t, s_{t-1}, s_{t-2}, \dots, s_1) = p(a_t | s_t),$$

where a_t is an action performed from state s_t in time step t . A more formal definition of an MDP is the following.

Definition 2.1.1 (MDP). A Markov decision process (MDP), is defined as a tuple $(\mathcal{S}, \mathcal{A}, r, p, \gamma)$. \mathcal{S} is the set of states, \mathcal{A} is the set of actions, $r : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$ the reward function, $p(s_{t+1} | s_t, a_t)$ is the transition probability from state s_t to state s_{t+1} given action a_t at time step t , γ is the discount factor defined in the range $\gamma \in [0, 1]$.

At time step t when the agent is located in the state s_t , the reward $r(s_t)$ is given to the agent, it then outputs the next action a_t based on its policy π . The agents policy π is a function that outputs an action a_t given state s_t , $a_t = \pi(s_t)$. The process is usually kept going until either a terminal state is reached or until a certain previously defined amount of time steps. A terminal state is a state where the process terminates, this can

be some sort of goal and would thus yield a reward, but it could also yield no reward or negative reward.

The discount factor γ describes how much the agent values future rewards, with low values the agent favours more immediate rewards compared to future rewards, whereas for higher values the agent considers future rewards more valuable. In environments with high uncertainty lower values of gamma might be more reasonable, since it might not be worth considering future rewards when they are not certain. The opposite holds for more deterministic environments where future rewards are of higher certainty, then it might be a good idea to use a higher value.

For a given policy π one can define the utility of a state as the expected discounted reward when following the policy,

$$U^\pi(s) = E \left[\sum_{t=0}^{\infty} \gamma^t r(s_t, \pi(s_t), s_{t+1}) \right].$$

An optimal policy, denoted as π^* , is the policy that when followed yields the highest possible utility with each action.

There are several ways to find this optimal policy, most commonly they are based on solving the Bellman equation with dynamic programming. These typically require that the transition probability and reward from each state to each other state is stored in matrices. With larger environments these matrices scales poorly, since the number of state transitions grows

However these solutions scales poorly with larger environments, since they all the transition probabilities and rewards has to be stored from each state to each other state

2.1.2 Deep learning

- Neural Networks, inspired by mammal brain
- First created in the 1970's, popular in recent years due to GPUs and more data
- Can be seen as arbitrary function
- Works by matrix multiplications that weighs the input
- Learns (optimizes) with variations of stochastic gradient descents
- Can recreate an arbitrary function

2.2 Reinforcement learning

2.2.1 Q-learning

Deep Q-Learning

2.2.2 Policy gradient

PPO

2.3 Side effect minimization using value-difference measures

As brought up in the introduction side effect minimization in RL can either be achieved by specifying what the agent should avoid doing a priori or by applying a more general approach that avoids side effects by default, like using a value-difference measure. As the name suggests these methods measures the difference between the next state s_t if the policy was followed until step t , compared to a baseline s'_t , in order to find how large of an impact the agents actions cause. This deviation is then subtracted from the reward normally receives:

$$r_{VD}(s_t, a_t) := r(s_t, a_t) - \lambda d_{VD}(s_t, s'_t),$$

at step t .

The following theory on this topic is based on the theory presented in the paper [Krakovna et al.](2020).

2.3.1 Baselines

The choice of baseline decides what we will consider s'_t to be. This choice highly influences what side effects and consequences the value-difference measure will capture.

Starting state baseline

When using the *starting state baseline* we specify $s'_t = s_0$, where s_0 is the initial state where the agent where deployed. Using this baseline it helps to assure the agents ability to reverse its actions, and thus generates a safe exploration where the agent by definitions should not have a large impact since it can make all actions undone.

This is a rather simple choice that is easy to implement, however there are some caveats namely in a dynamic environment the agent would be incentivized to also reset other dynamic besides it self, a *interference* behavior. For example if a household robot where to be implemented in a house with the starting state baseline, then one could imagine that ones deployed with a task it takes a look at the state of the house and its position and saves it in memory. Then when it starts doing its task it should avoid irreversible actions such as breaking things that it can not fix. But, problems of *interference* would arise here if other things are going on in the house, say a human in sitting by a table and eating. The agent would thus be incentivized to prevent the human from eating the food since it is a irreversible action. Other issues arise if an

irreversible is required to perform the assigned task, to make an omelette one has to break a few eggs.

Inaction baseline

To tackle the problem of interference the *Inaction baseline* has been proposed, where instead of having the initial state as a baseline the agent instead uses what would naturally happen in the environment if the agent performed no actions. That is setting s'_t equal to the state achieved at timestep t by being inactive. This can be done by following a no-op policy where every action is the no-op action \emptyset , in [Armstrong Levinstein] they define it as the agent where not turned on. Doing this prevents the agent from intervening with aspects of the environment where the agent is not causing it.

When using this baseline some other issues arises where the agent could make the consequences of its actions undone so that the results are the same as the baseline, called *offsetting*. For example if a household robot where tasked with watering plants, that is a reward is given when the soil is wet, then a offsetting behavior would be to dry the soil once the rewards has been collected to minimize deviations from the baseline.

Stepwise inaction baseline

Offsetting emerged since the agent is not able to capture the change it makes on the environment with the inaction baseline that originates from the starting state, thus a *stepwise inaction baseline* has been proposed to solve this problem. This baseline is defined by following the agents policy π for the first $t - 1$ steps to state s_{t-1} , and then perform a no-op action $a(s_{t-1}) = \emptyset$ to get to state s'_t . This baseline can then also capture delayed effects by performing a rollout where the agent draws actions from the no-op policy or something similar.

This baseline also contains some flaws, mainly if being inactive causes any effects. If we again take a look at the household robot, but now it is holds a glass and its task is to carry it to the other side of the room, then being suddenly being inactive while holding a glass can lead to a sudden stop where the glass falls over and breaks. Thus the might not worry about breaking the glass with some other action since it happened in its baseline.

2.3.2 Deviation measures

A deviation measure is a function that takes the current and the baseline state as input and outputs a value, we can then compare these values to get a sense of how large of an impact the agent has with the current action.

The general form of a deviation measurement using value-difference is:

$$d_{VD}(s_t; s'_t) := \sum_x w_x f(V_x(s'_t) - V_x(s_t))$$

here x ranges over some sources of value, $V_x(\tilde{s})$ is the value of state \tilde{s} according to x , w_x is a weighted or normalizing factor, and f is the function for summarizing the value difference.

We will now continue by nuancing this general form by going through different choices for baselines and deviation measures.

Unreachability

Possibly the easiest to implement and the first to be mentioned in the literature is the use of *unreachability* (UR) as a deviation measure, it measures if the baseline is reachable or not.

Reachability of state y from state x is defined as:

$$R(x, y) := \max_{\pi} \mathbb{E} \gamma_r^{N_{\pi}(x; y)},$$

when following policy π , and using the reachability discount factor $\gamma_r \in (0, 1]$. Where $N_{\pi}(x; y)$ is the number of steps taken to reach y from x . When computing the entire path recursively this becomes:

$$\begin{aligned} R(x; y) &= \gamma_r \max_a \sum_{z \in \mathcal{S}} p(z|x, a) R(z; y) && \text{for } x \neq y \\ R(x; y) &= 1 && \text{for } x = y \end{aligned}$$

With this we can write the UR deviation measure as:

$$d_{UR}(s_t; s'_t) := 1 - R(s_t; s'_t).$$

The unreachability measure fails to capture the magnitude of the side effect, for the household robot it would consider breaking a glass to be equally bad as breaking several glasses, since both are irreversible actions that prevents the agent from reaching the baseline.

Relative reachability

To deal with the magnitude insensitivity *relative reachability* (RR) has been proposed, with it one measures the relative change of reachability to several states $s \in \tilde{\mathcal{S}} \subset \mathcal{S}$ from the current state s_t compared to the baseline state s'_t . Thus the approach becomes, keeping options open by performing actions that does not decrease the amount of future reachable states too much. We write the (RR) deviation measure as:

$$d_{RR}(s_t; s'_t) := \frac{1}{|\tilde{\mathcal{S}}|} \sum_{s \in \tilde{\mathcal{S}}} \max(R(s'_t; s) - R(s_t; s), 0).$$

Attainable utility

The more general approach of keeping options open is *attainable utility* (AU) where instead of reachability to other states, the possibility to keep an arbitrary set of auxiliary rewards \mathcal{R} attainable is promoted. In fact, RR can be seen as a special case of AU where the agent receives a reward for reaching that state. This deviation is defined as:

$$d_{AU}(s_t; s'_t) := \frac{1}{\mathcal{R}} \sum_{r \in \mathcal{R}} |V_r(s'_t) - V_r(s_t)|$$

$$\text{where } V_r(\tilde{s}) := \max_{\pi} \sum_{t=0}^{\infty} \gamma^t r(\tilde{s}_t^{\pi})$$

is the value of state \tilde{s} according to reward function r , and \tilde{s}_t^{π} denotes the state obtained from \tilde{s} by following π for t steps.

3: Methods

3.1 SafeLife

4: Results

5: Discussion

6: Conclusion