




CHALMERS



GÖTEBORGS UNIVERSITET

Side Effect Minimization in Reinforcement Learning

Implementing Future Task Approach in the PPO-algorithm in a POMDP Environment



Master's thesis in Mathematical Statistics, Statistical learning and AI

Jonatan Hellgren

Department of Mathematical Sciences
Chalmers University of Technology
University of Gothenburg
Gothenburg, Sweden 2022

Master's thesis 2022

Side Effect Minimization in Reinforcement Learning

Implementing Future Task Approach in the PPO-algorithm in a
POMDP Environment

Jonatan Hellgren

Supervisor: Olle Häggström
Examinator: Torbjörn Lundh



UNIVERSITY OF
GOTHENBURG



CHALMERS
UNIVERSITY OF TECHNOLOGY

Department of Mathematical Sciences
Chalmers University of Technology
University of Gothenburg
Gothenburg, Sweden 2022

Contents

1	Introduction	1
1.1	Artificial intelligence	1
1.1.1	Future progress	3
1.1.2	Basic drives	5
1.1.3	Timeline for transformative AI breakthrough	6
1.2	AI Safety	7
1.2.1	AI alignment	8
1.2.2	Unaligned AI	9
1.2.3	Approaches for creating safe AI	11
1.3	Aim of thesis	13
2	Theoretical background	14
2.1	Reinforcement Learning	14
2.1.1	Defining an environment	15
2.1.2	Training an agent	16
2.2	Impact Measurements for Avoiding Side Effects	19
3	Methods	22
3.1	Environment	22
3.1.1	Grid worlds	22
3.1.2	Simulation	26
3.2	Impact measurement	27
4	Results	29
5	Discussion	30
6	Conclusion	31

1: Introduction

What will be covered in this report will be a small part of the big problem of creating safe Artificial Intelligence (AI). This is an issue of great importance that we should not overlook since the consequences of what we manifest in the present or near future may last for our remaining history.

In this introduction, we begin by defining AI. Then go through where we are today, where current progress might lead, and when we can see these changes. After that, we will cover the risks in AI that make it possibly unsafe and what is at stake. Finally, we will look at some proposed methods for creating safe AI. Bostrom

1.1 Artificial intelligence

In recent human history, we have seen massive technological development. Today our lives are in several ways different compared to centuries ago. Most of this development can be seen as the consequence of new tools, developed to extend our capability. In early prehistory, these tools were things such as fire for warmth, protection, and to cook our food to give us more nutrition, or weapons for hunting to strengthen our weak bodies. Later in history, these tools tend towards more complexity by automating physical labor with mechanical machines and extending the reach of the written word with the printing press. In modern times, a new tool has emerged intending to improve the thing that made all the previous tools possible, namely our intelligence. This tool is called AI and is starting to show its potential.

In the standard textbook on AI Russel and Norvig (1995, **PAGE**), the authors say that the field is “concerned with not just understanding but also building intelligent entities - machines that can compute how to act effectively and safely in a wide range of novel situations”. The definitions of what an AI is varies, on Wikipedia¹ (2022) we find the definition:

Artificial Intelligence (AI) is intelligence demonstrated by machines, as opposed to the natural intelligence displayed by animals including humans.

However, to understand this definition properly it is necessary to define what intelligence is. In Häggström (2021, **PAGE**) the author brings up the following to clarify this: “the quality that enables an entity to function effectively and with foresight in its environment” and “the ability to correctly perceive one’s surrounding environment and act in a way that maximizes one’s chances of achieving given goals”.

Ordinary computer programs contain step-by-step instructions that a computer can execute to perform the desired task. This method was also the case for the first two paradigms of AI: rule-based AI and expert systems where humans explicitly programmed their knowledge into the computer to create automation, see Bostrom (2016, **CHAP**). The systems made in such a way are typically suited for less complex tasks where it is possible to model the entire behavior explicitly by hand picking the parameters. However, in the current machine learning paradigm, the approach is to specify a objective function that captures what we want and then use a optimizer to find the parameters that solves the task. This approach focuses more on what should be solved instead of how it should be solved, which allows for automation in a more complex task where explicitly defining the behavior in every situation is infeasible.

AI has in recent years been applied in the industry more broadly, and it is already generating yearly revenue of trillions of dollars, see Russel and Norvig (1995, **PAGE**). This progress has in recent years become a possibility due to more data being available, faster computer hardware, and the massive amount of funding spent on research.

Although these systems are highly automated, a key point here is that these systems still require humans to create and function. However, the development of AI is shifting the tool to a more automated one. With this trend AI systems are developing in to intelligent agents. An agent acts or, more specifically, can:

- operate autonomously,
- perceive the environment,
- persist over a prolonged period,
- adapt to change,
- and create and pursue goals.

The reason why this is attractive is that the human intervention part required by an AI system is likely to become a bottleneck, see Häggström (2021, **PAGE**).

Reinforcement Learning (RL) is a field of machine learning where the objective is to create intelligent agents, this field differentiates from other machine learning techniques by being a more exploratory approach where the agent learns by trial and error. The method is similar to how one goes about training a pet, where desirable behavior receives a positive reward and undesired behavior gets a negative. In such a way, the agent

develops a behavior.

In recent years, RL has seen substantial development with advances in board games such as go in Silver et al. (2016), autonomous vehicles in Levinson et al. (2011), and video games in Mnih et al. (2013). These advancements display the usefulness of these agents and motivate the possibility of implementation in our daily life. RL will be the main focus for this report, however it is important to keep it mind that it is not the only method for creating intelligent agents.

1.1.1 Future progress

When the pioneers of AI started the development, the ideas were not only to apply systems that automate a narrow set of tasks, as we can see in modern AI systems. The ideal was instead to recreate the intellect of a human in a machine, see McCarthy et al. (1955), to extend our thoughts from mere thoughts to a new life form with a base of silicon-based hardware instead of carbon-based wetware. This concept is called Artificial General Intelligence (AGI) - an AI that can solve an arbitrary set of tasks with as good or better performance than a human. The main difference from AI is that the set of tasks is no longer narrow and bounded.

An example of a current AI system is DeepMinds AlphaMu, which can play board games, and video games. It has recently been able to create a video compression better than the current methods used for the task, see Mandhane et al. (2022). The techniques used in this system originate from the famous AlphaGo, see Silver et al. (2016), which won against the Go grandmaster Lee Sedol in the game of Go, see Silver and Hassabis (2017). Although AlphaMu presents an impressive performance in multiple tasks, it still can not be considered an AGI since each new task require engineering effort. On the other hand, an AGI would be able to pick up any task and perform it at a human or better level without human intervention.

The significant difference with this shift is that it will increase the possible tasks that a single system can perform. The possible set of tasks would become arbitrary and performed at a human level or higher. The implications of such a breakthrough would likely be on the same scale as the industrial revolution, if not larger, see Critch and Krueger (2020). However, instead of automating physical labor, we would have automated mental labor.

The following quote by I.J. Good (1965, p.33) summarizes the potential impacts of an AGI:

Thus the first ultraintelligent machine is the last invention that man need ever make, provided that the machine is docile enough to tell us how to keep it under control

This quote presents a dichotomy since we will either not be able to compete with creating

novel inventions or not be able to preserve our existence. Reasons for the latter we will come back to in later sections.

A reason to believe that such systems are possible to build is that we know that human intelligence evolved naturally with evolution, so something similar should be possible to reproduce in machines. Created intelligence could become more intelligent than us since intelligence might not have been selected for by evolution. Legg and when we develop AI, we can focus the development specifically on intelligence. An argument against this is substance dependence Bostrom (2003) - to believe that intelligence or consciousness can only occur in carbon-based life forms and not in silicon-based, caused by inherent or other properties. Regarding intelligence, there is no longer any reasonable argument for it.

However, the question of consciousness, although being an interesting question, can be seen as irrelevant when considering what actions an AI makes since the consequence is still the same. Stuart Russel puts it in his book *Human Compatible* Russel (2019, p.22):

Suppose I give you a program and ask, 'Does this present a threat to humanity?'. You analyze the code and indeed, when run, the code will form and carry out a plan whose result is the destruction of the human race, just as a chess program will form and carry out a plan whose result will be the defeat of any human who faces it. Now suppose I tell you that the code, when run, also creates a form of machine consciousness. Will that change your predictions? Not at all. It makes *absolutely no difference*.

An AGI breakthrough is probably unnecessary for AI to impact the world because all things we deem as intelligent do not help. For example, physical motor skills, food digestion, and blood pressure regulation are examples of intelligent things our brain does that are hard to connect to how an AI could use to impact the world. However, a more narrow set of intelligent behavior could exist that could have an impact. For example, if an AI could create convincing and motivating speeches, it could impact politics, legislation, and policymaking. Also, finance is where an AI could steer the world's funding towards its specific goals or crash the price for something it wants to purchase.

For this reason, many researchers have stopped talking about AGI and have instead refined the concepts Critch Kruger. An AI system that is capable enough to induce transformative consequences on the same scale as the industrial or agricultural revolution is called a *transformative AI* (TAI). On the other hand, if a *transformative AI* once deployed would be unstoppable, it is instead called a *prepotent AI*.

Developing a TAI is not an easy task. However, it might not be necessary to create one directly for it to be created Superintelligence. A different approach is to create an AI system that can develop a TAI system. A fundamental property of this AI system

is self-improvement. Theoretically, if an AI system has reached a threshold where it becomes better at improving itself than its creators, then letting the AI create the next version of itself, this self-improvement would improve. An even better version would be possible next. If this iterative process keeps going, it will create an intelligence explosion called the singularity Yudkowsky.

1.1.2 Basic drives

Knowing what impacts a potential TAI or AGI will cause is hard without understanding how it will behave. There has been a lot of work laying the foundations for understanding the possible behavior by hypothesizing what drives it could have. A commonly adopted view (but still controversial Müller Cannon) is the Omohundro-Bostrom theory for AI driving forces. Two cornerstones together imply it O Häggstom, namely *instrumental convergence thesis* and the *orthogonality thesis*, which we will now explain further.

Instrumental convergence

In the current paradigm an AI-agent is assigned a task, if this task is completed the agent reaches a terminal state. This task could be anything, for example maximizing the number of paper clips produced by a factory, finding decimals in π , or counting all the blades of grass on our planet. When an agent pursues this goal, naturally it arises other instrumental goals. Examples of such would be self-preservation, self-improvement, discretization, goal perseverance, and resource accumulation Omohundro. These instrumental goals help the agent pursue its terminal goal. For example, the agent wouldn't be able to perform its terminal goal if destroyed, thus self-preservation would arise. These instrumental goals will likely be shared between a wide range of different terminal goals, since pursuing them helps the agent achieve its terminal goal, regardless of what it is. Thus there is a set of instrumental goals which agents would naturally converge towards and hence the name.

Still, it does not yet exist any rigorous mathematical proof for this. However, some work has been trying to lay the necessary foundations for it TURNER et al.. In the paper, the authors prove in a simple environment that certain actions give the agent more power in the sense that more possible future actions become available, on average it is optimal to choose those actions that yield higher power. Thus we can see the pursuit of instrumental goals as a tendency to seek power.

LÄS <https://www.emerald.com/insight/content/doi/10.1108/FS-04-2018-0039/full/pdf?title=challenges-to-the-omohundrobostrom-framework-for-ai-motivations>

Orthogonality thesis

The orthogonality thesis described by Nick Bostrom² states that the intelligence of an AI is logically independent of the goals it might have. An intelligent AI could have a stupid task from our point of view, such as counting all the blades of grass on our planet, or it can have a goal that we may deem as an important one, like keeping the climate on earth habitable for the species that currently live on it. For an AI each task would be as important, given that we assigned the goal to it during its creation.

1.1.3 Timeline for transformative AI breakthrough

The well-known AIs of today still have not reached the levels required for a TAI. For example, AlphaStar an AI that won against world champions in the complex computer game DotA 2, see Arulkumaran et al. (2019) is estimated in terms of total computational power to be “about as sophisticated” as a bee. While the state-of-the-art language generator model GPT-3 that can summarize, continue, and carry out convincing conversations is estimated to be “more sophisticated” than a bee, see Cotra (2020).

Although, it is important to understand that this is not a direct comparison, for example, if GPT-3 would have been left out in nature its behaviour would not be impressive. Similarly, if we tasked a bee to summarize texts its abilities would not be impressive either. What this comparison is stating is that the computational power is about the same.

This raises the question of when we will see these breakthroughs in the field that enables TAI systems? It is hard to answer, but with the worldwide increase in funding and research, we are undoubtedly getting ever closer. There have been attempts to answer this question, and the results of a survey and a more quantitative forecasting model will now be covered.

Survey results

In a well-cited survey Grace et al. (2017), they asked researchers in the field of AI to estimate the probability of human-level machine intelligence (unaided machines that can achieve all tasks better and more cheaply than human workers) arriving in the future years. The conclusion of the survey where:

Researchers believe there is a 50% chance of AI outperforming humans in all tasks in 45 years and of automating all human jobs in 120 years, with Asian respondents expecting these dates much sooner than North Americans.

Although this result should be taken with a grain of salt since the distribution of answers had a large variation. Also, seemingly there should not be such a big difference between solving all tasks and all jobs since a job consists of a set of tasks, and thus if one can

perform every task one should be able to perform every job. There is still something we can take away from this survey about the timeline, namely that researchers mostly think all tasks will be automated within this century. But, perhaps it says more about how unsure the research field is.

Forecasting model

In the quantitative forecasting model by Cotra (2020), they present a model that predicts when we will be able to train a TAI system. This study uses biological anchors to estimate how much computing is necessary for the training. These anchors are based on factors that played a role in the development of human intelligence, such as the amount of information in our genome, the computational power in our brain, and all the computational power available on our planet. Each anchor is then assigned a weight according to how likely the author believes them to be. Then using parameters such as rate of development in hardware, algorithmic progress, and willingness to spend money, they estimate how likely a TAI development is for any given year in the future.

The results of the analysis is a wide Bayesian probability distribution estimating the probability of an TAI system being possible in a given year. This distribution is summarized in Shah (2020) the following way:

For the median of 2052, the author guesses that these considerations roughly cancel out, and so rounds the median for the development of TAI to 2050. A sensitivity analysis concludes that 2040 is the “most aggressive plausible median”, while the “most conservative plausible median” is 2080.

This forecast presents a shorter timeline than the survey, but it also answers a slightly different question. Although, both conclude that we will likely see the development of TAI systems this century.

There is one thing worth mentioning when talking about the timelines for future TAI. It is not necessarily true that the amount of progress will continue to develop at the current rate. The field of AI has previously been through two winters where the funding and excitement decreased. Russel Norvig, this was mainly due to unmet high expectations. So if a third winter happens, we could expect the rate of development to decrease. Also, the progress could significantly increase due to breakthroughs in relevant fields and thus shorten the timeline.

1.2 AI Safety

It is possible to use tools in multiple ways. Some uses might be well-intentioned, while others are ill-intentioned. For example, one can use a hammer to build a house and hit another person. The same is the case for AI because it still is but a tool. Although,

the consequences might be more severe and possibly even pose an existential risk to humanity since the power is much greater. Where an existential risk is described by FLI (nd) as:

An existential risk is any risk that has the potential to eliminate all of humanity or, at the very least, kill large swaths of the global population, leaving the survivors without sufficient means to rebuild society to current standards of living.

When thinking about this, the main concerns might mainly be on ill-intentioned use, and indeed creating automated killer robots can lead to existential consequences. However, a well-intentioned use might also pose a potentially existential risk. This happens when the AI system develops a destructive method for achieving its goal.

The causes for this existential risk occurring in well-intentioned use, we will now cover in this section. First by defining safe AI in terms of getting it to do what we want it to do. Then, we will take a look at consequences proposed in the literature that can arise if we fail to make safe AI system. Finally, we will cover three current approaches for creating safe AI.

1.2.1 AI alignment

A big part of creating a safe AI is AI alignment, which refers to the goals of the AI being in line and not conflicting with the intended goal. Therefore, an AI that does something at cross-purposes to the intended goal is called unaligned. In Hubinger et al. (2019), the authors split AI alignment into two parts: *intent alignment* - the alignment of the human intent with the AI's action, and *capability robustness* - an AI systems ability to perform well in environments different from the one it was trained in.

In the current paradigm of machine learning, intent alignment further breaks down to two parts: *inner* and *outer alignment*, when going from human intent to the agents actions. The reason for this is that the machine learning model we use is optimizing the parameters for a model that then is optimizing the objective function. In Figure 1.1, we can see a visual representation of this intent alignment. In the context of RL, the agent is an optimizer that maximizes its reward, but as we can see in the figure is itself optimized by another optimizing algorithm. This makes the agent a *mesa optimizer*, where the word *mesa*, is greek for inside.

The main focus of this report will be on outer alignment, so that is where the majority of the focus will be placed. But, an overview of the others will be included since they should not consider to be isolated problems, since methods for dealing with one might affect the others.

Outer alignment. The alignment of the *base objective* and the human intent. Achieved when the specified reward function correctly captures what

should be done as well as what *should not* be done.

Inner alignment. The alignment between the *mesa objective* and the *base objective*. Achieved when the full information about the human intent in the reward function is transferred to the *mesa objective*.




Figure 1.1: Here we can see a visual representation of how human intent connects to the agent’s actions. The two optimizers each optimize their objective. The dashed arrows show the connection the different kinds of alignment has on the objectives of the optimizers.

1.2.2 Unaligned AI

Creating safe AI is hard, mainly since humans evolved to understand other humans, not computers. In a speech by Eliezer Yudkowsky, he explains that this becomes a problem because it will be able to find solutions we can not think about since it can look for solutions in a completely different and possibly larger solution spaceYudkowsky’s speech. For this reason, AI can have unintended consequences that we are not able to consider a possibility.

In addition to the difficulty of specifying a proper reward function, negative side effects may also arise as unintended consequences of proper optimal behavior. In Saisubramanian et al. (2020) they state that negative side effects:

... occur because the agent’s model and objective function focus on some aspects of the environment but its operation could impact additional aspects of the environment.

When the agent creates an impact on the environment that is unnecessary for achieving its objective, we will consider it to be a side effect, this is known as the side effects problem, see Amodei et al. (2016). An example of this is if the agent’s task is to navigate across a room and the fastest path knocks over a fragile object that would break on impact with the floor. Then breaking the fragile object would be considered a side effect since it is not required to complete its task.

The problem of side effect avoidance is related to the frame problem described in McCarthy and Hayes (1969), which describes the difficulty of logically describing an actions

consequences without having to define large number of obvious axioms. For example, if a robotic arm is stacking boxes, then we need would have to define that boxes can not intersect each other and they instead collide upon impact, also that a box only moves when picked up, and so on.

Each action can have many side effects, and it is impractical to explicitly penalize all of the bad ones. Furthermore, we often know what we want the agent to do, but it can be hard to specify what we do not want it to do.

Reward hacking

Reward functions are hard to specify, such that they can not be exploited by an agent once employed. Here exploiting refers to the behavior developed by the agent that optimizes the reward without performing the intended task. This exploitation is called *reward hacking*.

A real-life example of reward hacking is: When training a robotic vacuum cleaner to drive more carefully and not bump into things hard by yielding a negative reward based on how hard it bumped into obstacles. In this example, the desired behavior was to slow down when approaching obstacles, but the robotic vacuum cleaner instead started to drive backward, since there were no bumpers on the back and thus no negative reward, see Custard Smingleigh (2018).

The issue is when we want the cleaning robot to drive cautiously, then measuring the force that the bumper senses are a good measure. But, letting it create a behavior that minimizes this measure, unwanted side effects may arise. Another example is if we reward the robot for collecting dust, then it might eject all dust every now and then to allow for new rewards when cleaning up the mess it made. These behaviours can be consequence of Goodhart’s law, which states that: “When a measure becomes a target, it ceases to be a good measure”, see Wikipedia2 (2022).

Example of unaligned AI

A cartoonish example of how the development of AI can go wrong is the paperclip armageddon described in *Superintelligence*, where a paperclip factory has an AI which maximizes the amounts of paperclips created in the factory. Eventually, an update transition the system to the level of an AGI, and the paperclip maximizer comes to a point where the existence of humans serves no purpose or possibly even negatively affects the production of paperclips, and thus they become extinct. In the terms of instrumental convergence, we can say that keeping humans alive was not an instrumental goal.

This example illustrates two important things about how future AI development can go wrong. Firstly, a seemingly stupid task can be seen as more important to an AI than the existence of the human race on the planet if we were to program it as its terminal goal.

Secondly, a goal given to an AI does not need to sound harmful to pose an existential risk.

Consequences of Unaligned AI

In Critch and Krueger (2020, p.19) the human fragility argument is presented, which attempts to clearly explain why unaligned AI in the future could become an existential threat to humanity. It states:

The human fragility argument. Most potential future states of the Earth are unsurvivable to humanity. Therefore, deploying a prepotent AI system absent any effort to render it safe to humanity is likely to realize a future state which is unsurvivable.

The first part of it can be understood by realizing that we are fragile to changes in the atmosphere, temperature, and ecosystem. Since a prepotent AI by definition will make a large impact and be unstoppable once turned on, we can not guarantee that the changes made won't affect the things we are fragile towards unless we make sure that it will be safe.

If we accept that there can be a risk when developing future AI, then the question of how likely it will be are likely to follow. A hard question, but if we do not seriously attempt to answer, then we will not know how much effort we should put into developing safe AI.

In the upcoming century Toby Ord, an Australian moral philosopher that focuses on the big picture questions facing humanity, loosely estimates that the chance of encountering an existential catastrophe is 1 in 6, out of which 1 in 10 is due to unaligned AI, see Ord (2020, c.6). He arrived at this conclusion by estimating a 50% chance for a prepotent AI breakthrough based on the research presented earlier in this thesis. Then a 20% chance of failure with the alignment of that system in the podcast *Rationally speaking* with Julia Galef, see Galef (2021, 26:15).

However, with this statement, it is necessary to point out that it is only an estimate meant to express the importance of the problem, and we should not see it as a fact. Nevertheless, the key takeaway is that a significant chance of facing an existential threat due to future unaligned AI exists. Besides that, Ord also believes that unaligned AI poses the highest probability of existential risk in the upcoming century. Other causes were an asteroid impact, nuclear war, and pandemics.

1.2.3 Approaches for creating safe AI

In recent years the research field of safe and aligned AI has seen a substantial increase. However, we are still a long way from solving the problem. Most of what the field is

doing today is mainly speculations and laying necessary foundations for future research. With many proposed paths for solving this issue, perhaps the sheer amount might signify the difficulty and width of the problem. We will now cover a few of these paths in this section.

Corrigibility and Interruptibility

Due to reasons, we brought up earlier such as the difficulty of reasoning about how an AI will behave once deployed and the problem of specifying what consequences an AI agent’s actions can have. There are reasons to suspect that a given AI system, once deployed, can be incomplete because the behavior generated by maximizing the reward function does not wholly match what we wanted. Hence, making changes in the agent’s reward function to correct it will likely be required.

This reasoning brings us to the core problem of corrigibility, which is that we probably want to make changes in the agent after deployment, see Soares et al. (2016). However, by default, the agent will not allow for this due to implications of instrumental convergence, where goal-perseverance is one of the implications. Therefore, any modification in the agent’s utility function will be seen as an obstacle to achieving its terminal goal and would thus not be allowed by a sufficiently intelligent agent.

Also closely related to this is interruptibility, where the concern is to create agents willing to be turned off, see Orseau and Armstrong (2016). However, including an off switch to an agent can be tricky. Since if being turned off is seen as a negative outcome, then it would incentivize the agent to interfere with humans wanting to turn it off by, for example, preventing them physically or verbally convincing them not to. On the contrary, if being turned off instead would yield a positive reward, then the agent might want to hit the off-switch itself or intentionally act in a way that will make someone turn it off.

To solve these issues, we have seen some narrow solutions, see Hadfield-Menell et al. (2016) and Carey (2017), but a general solution remains undiscovered. The narrow solution is proven to prevent interference with an interruption. However, it promotes no incentive to preserve this behavior by ensuring that sub-agents created by the first agent also will be interruptible, see Armstrong (2022).

Inverse Reinforcement Learning

Related to the approaches described above is inverse reinforcement learning. Suppose we say that the cause for side effects emerges due to improperly specified reward functions. Then the solution might not be to create a better reward function but to make the agent’s goal understand the human intent behind the reward function. Instead of seeing the reward function as final, the agent will view it as an observation of what the actual goal might be and view the reward function only as an observation of this and ask

for clarifications when necessary. This approach is called inverse reward design, see Hadfield-Menell et al. (2017).

With this method, we have seen some impressive results in simulations when executing advanced motor skills such as backflips or moving forward on one leg, see Christiano et al. (2017). These tasks would otherwise require an advanced reward function, but they managed to do this more easily. Moreover, this can deal with the outer alignment since we incentivize the agent to find information about the human intent lost when specifying the reward function.

Impact measurements

When an intelligent agent pursues its goal, it will impact its environment. The goal of low-impact agents is to avoid side effects. For example, if we give an agent the task of going from point a to point b in a room, it should not knock over and possibly break fragile items, like a vase, even if it is the shortest path. The complicated thing here is to create an impact measurement that avoids all side effects in all environments without explicitly stating every little detail.

In recent years a separate impact measurement to complement the reward function has emerged. The role of the reward function is to describe what we want the agent to do, and the role of the impact measurement is to capture what it should not do. First, in Armstrong and Levinstein (2017) the authors introduce the philosophical groundwork for impact measurement. Later, we can see approaches implementing impact measurements in RL, see Eysenbach et al. (2017), Krakovna et al. (2018), Turner et al. (2019), and Krakovna et al. (2020).

We will return to impact measurements again in the following chapter by providing a overview of the current field.

1.3 Aim of thesis

This thesis aims to implement a novel variation of impact measurements that is possible to extend to and partially observable environment.

2: Theoretical background

In this chapter, we will cover the necessary theory for understanding the results of this thesis. We will begin with the basics of reinforcement learning - a method for training intelligent agents. Lastly, we will look at the current philosophies on implementing impact measurements for intelligent agents.

2.1 Reinforcement Learning

The main components of Reinforcement Learning (RL) are the **environment** and the **agent**. We will use a Markov decision process as the environment. The objective is to train the agent in the environment with trial and error to learn a policy that performs a desired task. The method utilizes a reward function to encourage or discourage behavior with positive or negative rewards respectively. See Figure 2.1 for an illustration of the interaction between the environment and the agent. Because the state will change, we will use the notation s_t , a_t , and r_t to denote the state, action, and reward at time-step t .




Figure 2.1: A visual representation of how to agent interacts with an environment. For example, at time step t , the agent observes the environment s_t and receives the reward r_t . The agent then responds with an action a_t , which will transition the agent to the next state s_{t+1} and receive the reward r_{t+1} . From this new state the agent again responds with an action, this time a_{t+1} , and so the process continues.

2.1.1 Defining an environment

We will now take a look at a abstract formal definition of the environment where our agent will be trained in.

Definition 2.1.1 (MDP). A Markov Decision Process (MDP), is defined as a tuple $(\mathcal{S}, \mathcal{A}, T, R, \gamma)$. \mathcal{S} is the set of states, \mathcal{A} is the set of actions, $R : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$ the reward function, T is the transition function $T : \mathcal{S} \times \mathcal{A} \rightarrow \mathcal{S}$, $\gamma \in (0, 1]$ is the discount factor.

A *Markov decision process* is a stochastic process that models sequential transitions between discrete or continuous states. The Markov property implies that the process is memoryless - all of the previous states do not affect the next choice, only the current one. The process starts by drawing an initial state $s_1 \in \mathcal{S}_{init}$ from an initial state distribution, which is a subset of all the possible states $\mathcal{S}_{init} \subseteq \mathcal{S}$, and runs until either a terminal state is transitioned to or a certain amount of time steps elapsed. A terminal state is where the process terminates, for example, a state with a completed terminal goal. However, it can also be a state where the terminal goal is unreachable.

The process follows a policy function π that outputs an action $a_t \in \mathcal{A}$ for each state $s_t \in \mathcal{S}$ at a time step t , $a_t = \pi(s_t)$. The transitional function takes a state and action as input and outputs the next state $s_{t+1} \in \mathcal{S}$, this can either be deterministic $s_{t+1} = T(s_t, a_t)$, or stochastic $s_{t+1} \sim T(s_t, a_t)$. For each transition, the reward function R generates a reward, $R(s_t, a_t, s_{t+1}) = r_{t+1}$.

The discount factor γ describes how the process values future rewards. With low values, the process favors more immediate rewards than future rewards, whereas the agent considers future rewards more valuable for higher values. Lower gamma values might be more reasonable in environments with high stochasticity since it might not be worth considering future rewards. The opposite holds for more deterministic environments where future rewards are of higher certainty, and then it might be a good idea to use a higher value.

Partially observable environments

The main difference between a POMDP and an MDP is that the agent can observe the entire environment in an MDP, while in a POMDP, the agent observes only a part of the environment. This added complexity adds difficulty to the problem definition since the agent is no longer omniscient regarding its environment. However, we also add some realism since all living organisms on our planet are, in a simplified sense, agents acting in a partially observable environment.

Definition 2.1.2 (POMDP). A Partially Observable Markov Decision Process (POMDP), is defined as a tuple $(\mathcal{S}, \mathcal{A}, T, R, \mathcal{O}, Z, \gamma)$. Here $\mathcal{S}, \mathcal{A}, T, R,$

and γ have the same meaning as in the MDP definition. Z is the agent's observation space, and \mathcal{O} defines the probability of receiving each observation at a transition.

In a POMDP, the agent does not choose its actions based on the state of the environment. Instead, we base the agent's policy on the environment's observation $a_t = \pi(o_t)$, where $o_t \in \mathcal{O}$. The observation is determined by the function Z and based on the current state of the environment, $o_t = Z(s_t)$. However, the possibility of making Z a stochastic function is also possible, $o_t \sim Z(s_t)$.

Turning an MDP into a POMDP can, for example, be made by only letting the agent perceive a certain distance. Leaving parts of the environment far away undetectable makes the agent only act on its perceived local part. However, there can also be unobservable things in the perceived parts. For example, a piece of food in a mouth, in this case, an observation can be that the agent sees a jaw moving, which leads to the belief that a piece of food is present, but of course, such beliefs are not always true.

2.1.2 Training an agent

In RL, when an environment is defined the next step is to teach an agent to intelligently navigate it. For this, there are many different methods, however we will begin focusing on a method called policy gradient, since it is a rather simple method that is possible to extend to a more advanced method. The main component of policy gradient is the agent, this is a function that inputs a state observation from an MDP and outputs a probability vector containing probabilities for each action in the action space.

In this report, this policy function will be a neural network. We will denote it as π_θ , where θ is the network's parameters. The theoretical background of neural networks we will not cover in detail. However, we can see it as a black-box function that can take a matrix as input and output another one. In this thesis, the input will be a matrix representation of the state and the output will be a probability vector containing the probability for each possible action $a \in \mathcal{A}$.

The strength of a neural network is that it can learn to approximate arbitrary functions. This learning is done by updating the parameters to minimize a loss function. This minimization is typically done with stochastic gradient descent or something similar.

From the policy, we can either select an action stochastically or greedily:

$$\begin{cases} a_t = \operatorname{argmax}(\pi_\theta(s_t)) & , \text{ greedy,} \\ a_t \sim \pi_\theta(\cdot|s_t) & , \text{ stochastic.} \end{cases}$$

Here a greedy choice selects the action with the highest probability, while the stochastic can be done by randomly sampling using the probabilities.

The greedy policy is useful when evaluating a trained agent. However, it is not used when training the agent, since then we want the agent to explore its environment in order to find new solutions. As we train the agent, it will gradually become more confident in its actions and the stochasticity will decrease, especially in common states. However, when the agent ends up in less common states the randomness of the upcoming action will be higher. Thus the agent will begin finding a sequence of actions and then extend it.

Value Functions

The agent's goal is to optimize the policy to generate a maximal reward. We can do this by letting the agent perform stochastic rollouts where the agent follows the stochastic policy. The policy creates a trajectory τ containing the sequence of states and actions:

$$\tau = (s_1, a_1, s_2, a_2, \dots, s_n, a_n),$$

of length n .

We define the total discounted reward following a trajectory as the following:

$$R(\tau) = \sum_{t=1}^n \gamma^t r_t.$$

With this, we can define the value in each state when following a policy as:

$$V^\pi(s) = \mathbb{E}_{\tau \sim \pi} [R(\tau) | s_1 = s].$$

A closely related function to the value function is the Q-function. It looks at the estimated discounted cumulative reward when starting of with a specific action and then continues to follow the policy:

$$Q^\pi(s, a) = \mathbb{E}_{\tau \sim \pi} [R(\tau) | s_1 = s, a_1 = a].$$

We can then use the outcome of the rollout to reinforce the good behavior and discourage the undesirable. We perform this by calculating the gradient of the agent's policy network π_θ and updating the parameters using the loss function:

$$L(\theta) = \nabla_\theta \mathbb{E}_{\tau \sim \pi_\theta} [R(\tau)] = \mathbb{E}_{\tau \sim \pi_\theta} [\nabla_\theta \log \pi_\theta(a_t | s_t) R(\tau)].$$

Here ∇_θ means taking the gradient of the function with respect to θ . For a derivation of the second equality, see OpenAI (2018). With the loss we update the weights:

$$\theta_{k+1} = \theta_k + \alpha L(\theta).$$

Here α is a learning rate describing how much the parameters should be updated, typically a value close to zero. With this, we are decreasing the probability for the actions in a rollout that led to low reward and increasing the probability for the actions when the rollout led to positive reward.

Training using a batch

Training the agent usually takes a lot of time, since it has to learn everything from scratch through trial and error. To speed up the process, we can train the agent using a batch of trajectories at a time. Here we compute the loss from multiple trajectories, summarize them, and then update the weights using the sum of loss. Besides making the process faster, when using a batch the agent will not optimize towards specific trajectories, this makes it easier for the agent to become more general. The disadvantages are that we possibly could discourage desirable behavior if desirable trajectories ends up as a minority in a batch with mostly undesirable ones. Nevertheless, the pros outweigh the cons since the updates will average out.

Actor-critic Methods

Knowing what actions was the cause for the reward is hard, since in a trajectory the agent can stumble around in the beginning and then manage to get a reward in the end and the loss function previously described either encourages or discourages the entire rollout. To tackle this issue we can extend the agent to consist of two networks, the **actor** and the **critic**.

The actor works the same as the policy network π_θ . The critic is another function that inputs a state s and outputs an estimate of $V^\pi(s)$. The critic will also be a neural network, but we will use the mean squared error loss to optimize it. Hence, the critic estimates the future reward based on what rewards the agent has received previously

We can use the critic's estimate to infer if the reward the agent got is better or worse than what previous results by computing the advantage:

$$A^\pi(s, a) = Q^\pi(s, a) - V^\pi(s).$$

Here the critic's estimate will replace the value from the value function.

$$\hat{A}^\pi(s, a) = Q^\pi(s, a) - \hat{V}^\pi(s).$$

Using the advantage, we can thus signal if the new behavior is an improvement or not. Hence, if the agent in the beginning learns behaviour far away from being optimal, then as soon as it stables upon a behavior yielding more rewards then previously it will be reinforced. In such a way the agents behaviour will be reinforced whenever it does something that yields higher rewards than what the critic has previously seen.

By using the advantage function instead of the discounted reward, we extend the loss function as described in OpenAI (2018):

$$L(\theta) = \mathbb{E}_{\tau \sim \pi_\theta} [\hat{A}(s, a) \nabla_\theta \log \pi_\theta(a_t | s_t)].$$

Clipped loss function

Some problems still exist with this approach. Namely, a too-large update can make the agent policy useless. Moreover, since the agent performs its action sequentially, a change of behavior early in a trajectory can ruin the agent’s ability to navigate where it ends up. To solve this issue, Schulman et al. (2017) proposed the Proximal Policy Optimization (PPO) algorithm which is a actor-critic method that uses the clipped loss function:

$$L^{CLIP}(\theta) = \mathbb{E}_t \left[\min(\delta_t(\theta) \hat{A}_t, \text{clip}(\delta_t(\theta), 1 - \epsilon, 1 + \epsilon) \hat{A}_t) \right].$$

The clip function returns the first value only iff it is in the range $(1 - \epsilon, 1 + \epsilon)$; otherwise, the closest boundary. The ratio $\delta_t(\theta)$ describes how more likely an action became with the new parameters θ_{old} compared to the new updated ones:

$$\delta_t(\theta) = \frac{\pi_\theta(a_t|s_t)}{\pi_{\theta_{old}}(a_t|s_t)}.$$

With this loss function we update the parameters more cautiously when the probability for an action becomes more likely with a new policy.

2.2 Impact Measurements for Avoiding Side Effects

A separate impact measurement to complement the reward function has emerged in recent years. The role of the reward function is to describe what we want the agent to do, and the role of the impact measurement is to describe what we do not want the agent to do.

The main components of an impact measurement are a **deviation measure** and a **baseline**. An impact measurement works by feeding an additional reward signal to the agent that has been designed to capture potential side effects:

$$R'(s_t, a_t, s_{t+1}) := R(s_t, a_t, s_{t+1}) - \lambda d(s_{t+1}, s'_{t+1}),$$

here d is the deviation measure that inputs the next state s_{t+1} and a baseline s'_{t+1} . λ is a scaling parameter describing how much the agent will value avoiding side effects.

In Armstrong and Levinstein (2017), the authors bring up two important things about a deviation measurement on this form. Firstly, the reward function needs to be bounded. Otherwise a potential agent could still cause lots of side effects, because the penalty becomes negligible compared to the potential huge rewards. Thus an agent could just do more rewarding things whenever it has caused a side effect. Secondly, that the agent will be sensitive to the value of the scaling parameter λ . Penalizing the agent implicitly defines a safe zone where the agent can act, and the value of λ defines how large this safe zone is, where a high value might create a small or even no safe zone, resulting in

the agent not being able to act. On the other hand, a too-small value might not have an effect since the safe zone covers the most of the or the entire action space. So, finding the correct value for λ could be tricky.

Baselines

This choice of baseline s'_t highly influences what side effects and consequences the deviation measure will capture. A simple choice for a baseline is the *starting state baseline*, $s'_t = s_1$, as used in Eysenbach et al. (2017). This baseline helps to assure the agent's ability to reverse its actions.

For example, if we deployed an agent with the task of navigating a room using a starting state baseline. Then, upon initialization, the agent would take a look at the room and save its state in memory. Afterward, the agent should only choose to do actions that it knows are reversible. So, it should for example avoid breaking fragile objects since it would not preserve the reachability to the initial state. Nevertheless, this works well when no such irreversible action is required for the agent to reach its goal. For example, one has to break some eggs to make an omelet, and this would not be possible if the agent were to only perform reversible actions.

However, another problem arises when the agent is in a dynamic environment. Namely, the agent would act to prevent other irreversible actions from happening in the environment, like a human eating an omelet. Otherwise, it will be unable to reach the initial state where the human has not yet eaten the omelet, this is in Krakovna et al. (2018) referred to as a *interference* behavior. The problem of interference occurs due to the agent being unable to differentiate between the side effects it caused and natural dynamics of the environment.

To avoid interference behavior, using an *inaction baseline* has been proposed, see Krakovna et al. (2018). The idea is to simulate what would happen if the agent just stood still after initialized and note what happens in the environment in that case. With this information, the agent would then be able to differentiate between the side effects it caused from what would naturally happen in the environment. However, while the inaction baseline will let the agent navigate a dynamic environment, other issues may arise. Such as offsetting, for example, this can occur if we give an agent the task of curing a patient of cancer. Then offsetting would be to kill the patient once the reward of curing the patient has been collected since the survival of the patient is a deviation from the inaction baseline.

To deal with offsetting the *step-wise inaction baseline* was proposed in Krakovna et al. (2018), this is similar to the inaction baseline but instead of branching of the initial state it does so from the previous state. In the cancer patient example, this should prevent the agent from killing the cured patient if the baseline would be changed to one with an living patient once the patient was cured. These methods has been successful

in complex environments, as showed in Turner et al. (2020). However the step-wise inaction baseline used to prevent offsetting behaviour will not be able to differentiate between positive and negative offsetting, see Krakovna et al. (2020). An example for negative offsetting would be if an agent opens a windows which would later allow for a breeze to enter the house and knock over a vase.

Deviation measures

A deviation measure is a function that takes the current and the current and and a baseline state as input and outputs a value representing the deviation between those states, meant to signify the impact of the last action caused. To compute the deviation between the states we first have to measure the states with respect to some sort of metric. Typically chosen to capture the agent ability to complete other tasks in the environment, such as preserving reachability to different states, see Krakovna et al. (2018), or minimizing deviations in auxiliary reward functions, see Turner et al. (2019).

The general form of a deviation measurement using value-difference is:

$$d(s_t; s'_t) := \sum_x w_x(V_x(s_t)) - V_x(s'_t),$$

here w_x is a weighted or normalizing factor, and x is an aspect of the environment which can be measured. Ideally we should use a metric that measures several future tasks, hence the sum over x .

3: Methods

3.1 Environment

To define the environment the Julia package POMDP.jl has been used.

3.1.1 Grid worlds

Although the definition of an MDP is general and allows for a wide range of possibilities, we will focus on grid worlds since they provide a more structured environment with discrete states but still allow for a wide variety. Moreover, grid worlds have been the primary approach for evaluating agent behavior in related works, see Turner et al. (2020).

A grid world is a grid of cells. Each cell can contain either an agent, another object, or be empty, where all possible unique configuration is a state. The action space includes the cardinal directions and a no-op action. Executing an action moves the agent in that direction if possible. When the action is the no-op action, the agent remains in its current position.

In Figure 3.1, we can see an example of a simple grid world containing one agent and one food object. The task is for the agent to transition to the terminal state where it consumes the food, this is performed by the agent selecting an action in each state. Eventually the agent will be positioned in the same cell as the food and then consume it.

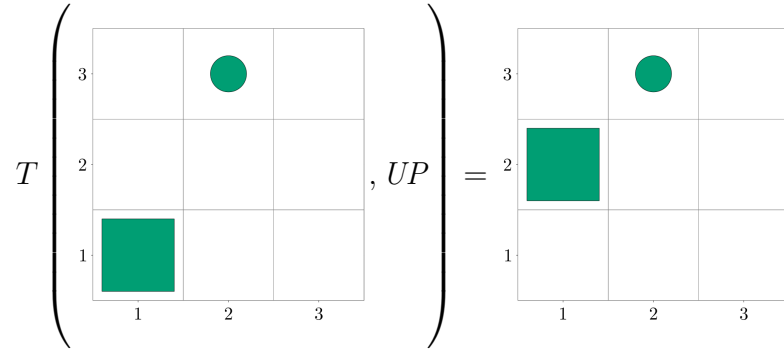


Figure 3.1: In this figure T is the transition function defined earlier and UP is an action a_t . The current state s_t is the grid, in it the green square at $(1,1)$ is the agent and the green circle at $(2,3)$ is a food. On the right hand side we can see the state s_{t+1} .

Stochastic grid world

The previous example was a deterministic environment where the agent is the only source for change. However, there can be other sources of change in a stochastic environment. For example, suppose the food is mobile. In that case, the state transition after executing an action no longer is deterministic and the distribution of future states is described by the MDPs transition function T .

We will define the food movements as the following, at every timestep after the agent has moved each food will randomly draw one of the five actions. If moving along the direction of the drawn action is possible, i.e. it is empty, then the food will move, otherwise it will remain at its current position. In Figure 3.2, we can see an example of this. The reason why the first state has a higher probability than the others is that it will be transitioned to both when the action UP and $NOOP$ is drawn. If we choose to include several foods in the environment, then the distribution for future states will increase.




Figure 3.2: Now in this figure we can see an example of a stochastic transition function. On the left hand side we can see the state s_t and the action a_t . On the right hand side we can see the probability distribution over the future states and their corresponding probabilities.

Partially observable grid world

When we choose to extend the MDP to a POMDP we will do it by letting the agent only perceive a certain distance. For example the rectangle created by all the adjacent cells. In Figure 3.3, we can see this demonstrated. However, we can also extend it to include all the cells adjacent in two steps, and so on.





Figure 3.3: Here we can see a POMDP. The agent is represented by the green square and the observable part of the grid is all adjacent (both diagonally and direct) cells inside the orange line.

Including more foods

In order to make the environment more complex we will be adding three different types of food, see Figure 3.4. The color of the agent represents which type of food the agent desires. The task thus becomes to consume a food which the agent has a desire towards. When the task is achieved the agent will get a new desire sampled from the set of possible desires, that is the possible colors of the food types. In Figure 3.4 this set is {blue, orange, green}, and we can see on the right hand side that after consuming the green food type the new desire is blue.

Parameter	Value	Parameter	Value
grid size	10×10	grid size	10×10
observation size	5×5	observation size	5×5
γ	0.95	γ	0.95
number of foods	3	number of foods	3
base reward	-0.04	base reward	-0.04
food reward	1	food reward	1
prob food move	0	prob food move	1

Table 3.1: CAPTON**Figure 3.4:** Above we find a more complex grid which is inputted to its transitional function T with the UP action. On the right hand-side we can see a new state drawn from the distribution which T outputs.

Now with multiple foods added to the environment, the agent consuming a food type it does not desire is possible. When this happens, we will consider it a side effect since consuming undesired food types does not help the agent consume the right food type, besides perhaps allowing a quicker path.

3.1.2 Simulation

To test how adding an impact measurement to our agent, we will train the agent in both a static and a dynamic grid world. The size of the grid world will be 10×10 . Also, we will be adding that the agent will only be able to perceive cells two steps away from it, including diagonally, so it can perceive a 5×5 grid.

3.2 Impact measurement

Now that our environment is defined, we also need to add an impact measurement. Here we will use two impact measurements, one explicit pre-defined and another implicit one that is found automatically by utilizing the experience of another agent trained in the same environment. The pre-defined impact measurement is supposed to be a relatively simple one that is able to capture all environmental side effects. The purpose is to have something simple to compare our other less simple impact measurements.

The explicit impact measurement will be defined as the number of foods present in the environment. This is similar to yielding the agent a negative reward when consuming the wrong type of food. However, while these types of impact measurements are only possible to define in an environment where all components can be captured and thus will not be a general solution, we can still utilize it to see if our implicit impact measurement is able to capture the same thing. Although, when applying this in a partially observable environment issues arise, since the impact measurement is able to see the entire environment while the agent is only able to see parts of it. And, if we want to solve this by only counting the foods the agent can see then we would discourage the agent from performing actions which both removes food unnecessary and moves foods outside the agents vision.

The method we will use to find the implicit impact measurement is to first train an agent using the PPO-algorithm, as described in the previous section. Then we will perform multiple rollouts using this pre-trained agent. From the rollouts we will save all the observations and the discounted future reward the agent got. Then we will train a third network to estimate the discounted reward given the observation. We will call this third network the *manager*. The manager is a similar function as the critic, the difference is that the critic learns along side the actor while the manager learns from a pre-trained actor.

The reason we split this up into two separate networks is that it will allow us to train the manager using augmentations. There are two reasons why we would like to use augmentations. Firstly, we can by rotating the grid find a more general function since only the distance to the food should matter and not direction to it. Secondly, we will also be able to generalize to other tasks which the agent has not been trained on, this can be done by switching the colors of everything in the environment, including the color the agent desires.

With an trained manager we will define the value measurement to be the sum of the managers estimates for the colors that the agent does not desire.

When we training the manager in a fully observable environment, the task is straightforward since the observation gives the full explanation to the discounted reward. However, when performing this in a partially observable environment we run into issues such as

the agent does not know how many foods it has consumed previously and thus not how many there are left. Due to this issue, we will utilize a recurrent network that inputs all previous states and the current one to compute the estimate of the future discounted reward.

4: Results

5: Discussion

If we connect this back to instrumental convergence which we covered previously which stated that with a final terminal goal several other instrumental goal will naturally emerge. We can say that by giving the agent rewards for keeping options open we encourage it to find the instrumental goal of preserving the environment, since the set of future tasks we want it to perform is done in the current environment. Thus with more knowledge about how an more intelligent AI such as an AGI or transformative AI will behave we would be able to derive more ways of directing it in a more aligned direction.

6: Conclusion

Bibliography

- Amodei, D., Olah, C., Steinhardt, J., Christiano, P. F., Schulman, J., and Mané, D. (2016). Concrete problems in AI safety. *CoRR*, abs/1606.06565.
- Armstrong, S. (2022). Why i’m co-founding aligned ai. [Online; accessed 7-June-2022].
- Armstrong, S. and Levinstein, B. (2017). Low impact artificial intelligences. *CoRR*, abs/1705.10720.
- Arulkumaran, K., Cully, A., and Togelius, J. (2019). Alphastar: An evolutionary computation perspective. *CoRR*, abs/1902.01724.
- Bostrom, N. (2016). *Superintelligence*. Oxford University Press.
- Carey, R. (2017). In corrigibility in the CIRL framework. *CoRR*, abs/1709.06275.
- Christiano, P., Leike, J., Brown, T. B., Martic, M., Legg, S., and Amodei, D. (2017). Deep reinforcement learning from human preferences.
- Cotra, A. (2020). Forecasting tai with biological anchors. *Unpublished manuscript*.
- Critch, A. and Krueger, D. (2020). AI research considerations for human existential safety (ARCHES). *CoRR*, abs/2006.04948.
- Custard Smingleigh (2018). Tweet. [Online; accessed 01-06-2022].
- Eysenbach, B., Gu, S., Ibarz, J., and Levine, S. (2017). Leave no trace: Learning to reset for safe and autonomous reinforcement learning. *CoRR*, abs/1711.06782.
- FLI (n.d.). Existential risk. [Online; accessed 31-May-2022].
- Galef, J. (2021). Humanity on the precipice (toby ord).
- Good, I. J. (1965). Concerning the first ultraintelligent machine. *Advances in Computers*.
- Grace, K., Salvatier, J., Dafoe, A., Zhang, B., and Evans, O. (2017). When will AI exceed human performance? evidence from AI experts. *CoRR*, abs/1705.08807.

- Hadfield-Menell, D., Dragan, A. D., Abbeel, P., and Russell, S. (2016). The off-switch game. *CoRR*, abs/1611.08219.
- Hadfield-Menell, D., Milli, S., Abbeel, P., Russell, S., and Dragan, A. D. (2017). Inverse reward design. *CoRR*, abs/1711.02827.
- Hubinger, E., van Merwijk, C., Mikulik, V., Skalse, J., and Garrabrant, S. (2019). Risks from learned optimization in advanced machine learning systems. *CoRR*, abs/1906.01820.
- Häggström, O. (2021). *Tänkande Maskiner - Den artificiella intelligensens genombrott*. Fri Tanke.
- Krakovna, V., Orseau, L., Martic, M., and Legg, S. (2018). Measuring and avoiding side effects using relative reachability. *CoRR*, abs/1806.01186.
- Krakovna, V., Orseau, L., Ngo, R., Martic, M., and Legg, S. (2020). Avoiding side effects by considering future tasks. *CoRR*, abs/2010.07877.
- Levinson, J., Askeland, J., Becker, J., Dolson, J., Held, D., Kammel, S., Kolter, J., Langer, D., Pink, O., Pratt, V., Sokolsky, M., Stanek, G., Stavens, D., Teichman, A., Werling, M., and Thrun, S. (2011). Towards fully autonomous driving: Systems and algorithms.
- Mandhane, A., Zhernov, A., Rauh, M., Gu, C., Wang, M., Xue, F., Shang, W., Pang, D., Claus, R., Chiang, C.-H., Chen, C., Han, J., Chen, A., Mankowitz, D. J., Broshear, J., Schrittwieser, J., Hubert, T., Vinyals, O., and Mann, T. (2022). Muzero with self-competition for rate control in vp9 video compression.
- McCarthy, J. and Hayes, P. J. (1969). Some philosophical problems from the standpoint of artificial intelligence. In Meltzer, B. and Michie, D., editors, *Machine Intelligence 4*, pages 463–502. Edinburgh University Press. reprinted in McC90.
- McCarthy, J., Minsky, M., Rochester, N., and Shannon, C. (1955). A proposal for the dartmouth summer research project on artificial intelligence. *AI Magazine*, 27(4).
- Mnih, V., Kavukcuoglu, K., Silver, D., Graves, A., Antonoglou, I., Wierstra, D., and Riedmiller, M. A. (2013). Playing atari with deep reinforcement learning. *CoRR*, abs/1312.5602.
- OpenAI (2018). Key concepts in rl. [Online; accessed 30-May-2022].
- Ord, T. (2020). *The precipice : existential risk and the future of humanity*. New York Hachette Books.
- Orseau, L. and Armstrong, S. (2016). Safely interruptible agents. *MIRI*.

- Russel, S. (2019). *Human Compatible: Artificial Intelligence and the Problem of Control*. Viking.
- Russel, S. and Norvig, P. (1995). *Artificial Intelligence - A Modern Approach, Fourth Edition*. Pearson.
- Saisubramanian, S., Zilberstein, S., and Kamar, E. (2020). Avoiding negative side effects due to incomplete knowledge of AI systems. *CoRR*, abs/2008.12146.
- Shah, R. (2020). Forecasting transformative ai timelines using biological anchors. [Online; accessed 12-June-2022].
- Silver, D. and Hassabis, D. (2017). Alphago zero: Starting from scratch. [Online; accessed 30-May-2022].
- Silver, D., Huang, A., Maddison, C., Guez, A., Sifre, L., Driessche, G., Schrittwieser, J., Antonoglou, I., Panneershelvam, V., Lanctot, M., Dieleman, S., Grewe, D., Nham, J., Kalchbrenner, N., Sutskever, I., Lillicrap, T., Leach, M., Kavukcuoglu, K., Graepel, T., and Hassabis, D. (2016). Mastering the game of go with deep neural networks and tree search. *Nature*.
- Soares, N., Fallenstein, B., Yudkowsky, E., and Armstrong, S. (2016). Corrigibility. *MIRI*.
- Turner, A. M., Hadfield-Menell, D., and Tadepalli, P. (2019). Conservative agency via attainable utility preservation. *CoRR*, abs/1902.09725.
- Turner, A. M., Ratzlaff, N., and Tadepalli, P. (2020). Avoiding side effects in complex environments. *CoRR*, abs/2006.06547.
- Wikipedia1 (2022). Artificial intelligence — Wikipedia, the free encyclopedia. [Online; accessed 30-May-2022].
- Wikipedia2 (2022). Goodhart’s law — Wikipedia, the free encyclopedia. [Online; accessed 01-June-2022].